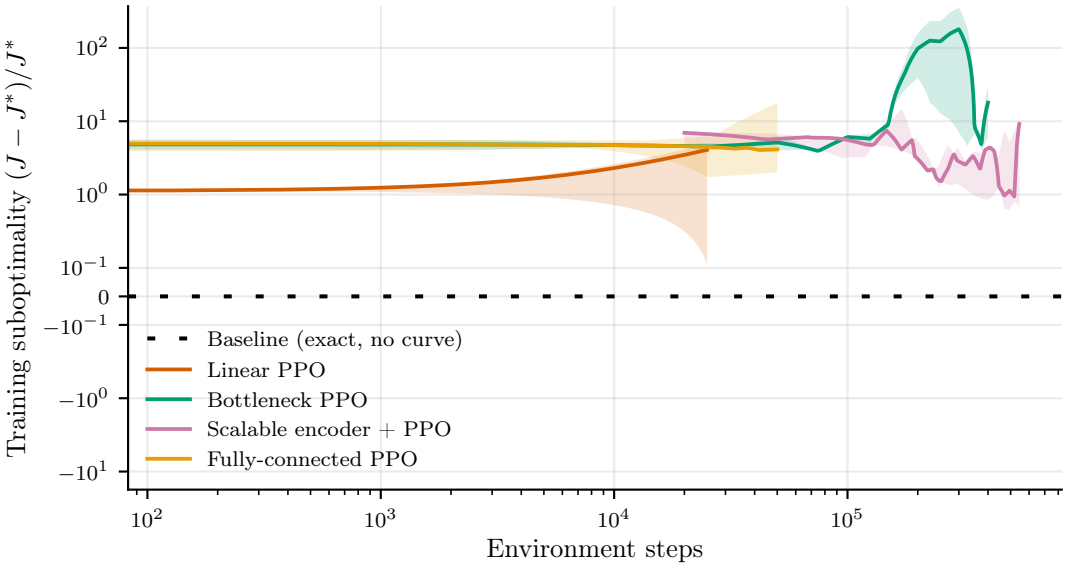

Training suboptimality $(J - J^*)/J^*$
10^2
10^1
10^0
10^{-1}
0
-10^{-1}
-10^0
-10^1
Baseline (exact, no curve)
Linear PPO
Bottleneck PPO
Scalable encoder + PPO
Fully-connected PPO
10^2
10^3
10^4
10^5
Environment steps